

SAGNNIK BISWAS

+91-8158867604

sagnnikbiswas2002@gmail.com

Sagnnik Biswas

Sagnnik

EDUCATION

Manipal Institute of Technology

BTech - Electronics and Communication (Data Science minor) - CGPA - 8.34

Oct 2020 – Nov 2024

Manipal, Karnataka, India

EXPERIENCE

Research Internship

Rakuten India

Jan 2024 - Feb 2025

Bangalore

- Designed a cross-modal GNN–LLM architecture for structured banner layout generation.
- Built spatial layout GNNs and aligned graph embeddings with LLM representations to improve multimodal reasoning.
- Scaled end-to-end training using multi-GPU DDP and DeepSpeed ZeRO-3, enabling efficient training and inference at scale.

Research Internship

IIT Kharagpur - Under the guidance of Prof. Somak Aditya

Jan 2023 – Dec 2023

- Constructed entity–event causal scene graphs from collision datasets to represent spatio-temporal interactions
- Designed counterfactual QA tasks over these graphs to evaluate and enhance LLM reasoning capabilities.

Robotics Internship

Aeonix Research and Innovations LLP

Jul 2023

Kolkata

- Built an automated industrial metrology system using YOLOv8 + SAM + geometric CV pipelines.
- Reduced inspection time from 15–20 min to less than 1 min.
- **Awarded Best Paper** – RCAAII 2023 (Springer LNEE).

RoboManipal

Electronics Subsystem Member

Dec 2020 – Dec 2023

MIT, Manipal

- Developing and Designing Circuits, CADs.
- Automation of Robots with computer vision.
- Selected for Nationals in ABU Robocon, 2022 and 2023.
- First place in Circuit Fixer, an event organized by COEP, Pune.

PROJECTS

Docker MCP Bridge — LLM Tool Orchestration Infrastructure ↗

Dec 2025

- Built a secure, container-native middleware enabling dynamic discovery, generation, and sandboxed execution of LLM tools for automated workflows via MCP.
- Implemented Redis-backed tool registry, user-scoped MCP sessions, and interrupt/resume workflows.
- Orchestrated isolated tool execution using Docker with FastAPI streaming APIs.

Multi-Agent Web Researcher ↗

Nov 2025

- Designed a LangGraph-based stateful multi-agent research engine with Redis checkpointing and resumable execution.
- Implemented real-time SSE streaming, citation synthesis, and multi-source search pipelines.
- Built a Streamlit frontend for interactive research workflows.

Ollama ScaleWright — GPU LLM Serving Platform ↗

Nov 2025

- Built a GPU-aware LLM serving platform with Celery-based async inference and Redis-backed model tracking
- Implemented a knapsack-based VRAM eviction algorithm for dynamic model loading/unloading.
- Enabled real-time streaming inference for low-latency, multi-user LLM workloads.

TECHNICAL SKILLS

Developer Tools / Frameworks: PyTorch, Transformers, LangGraph, FastAPI, Docker, Redis, Celery

Additional: Streamlit, React, MongoDB, OpenCV

EXTRACURRICULAR

- Co-presenter for OpenCV workshop
- In the organising team of Robowars, a premier event in TechTatva'23

Aug 2022
Oct 2023

CERTIFICATIONS

- Data Science Specialization: JHU - Coursera